

Classification and Regression Trees

Trees and Rules

Goal: Classify or predict an outcome based on a set of predictors

The output is a set of **rules**

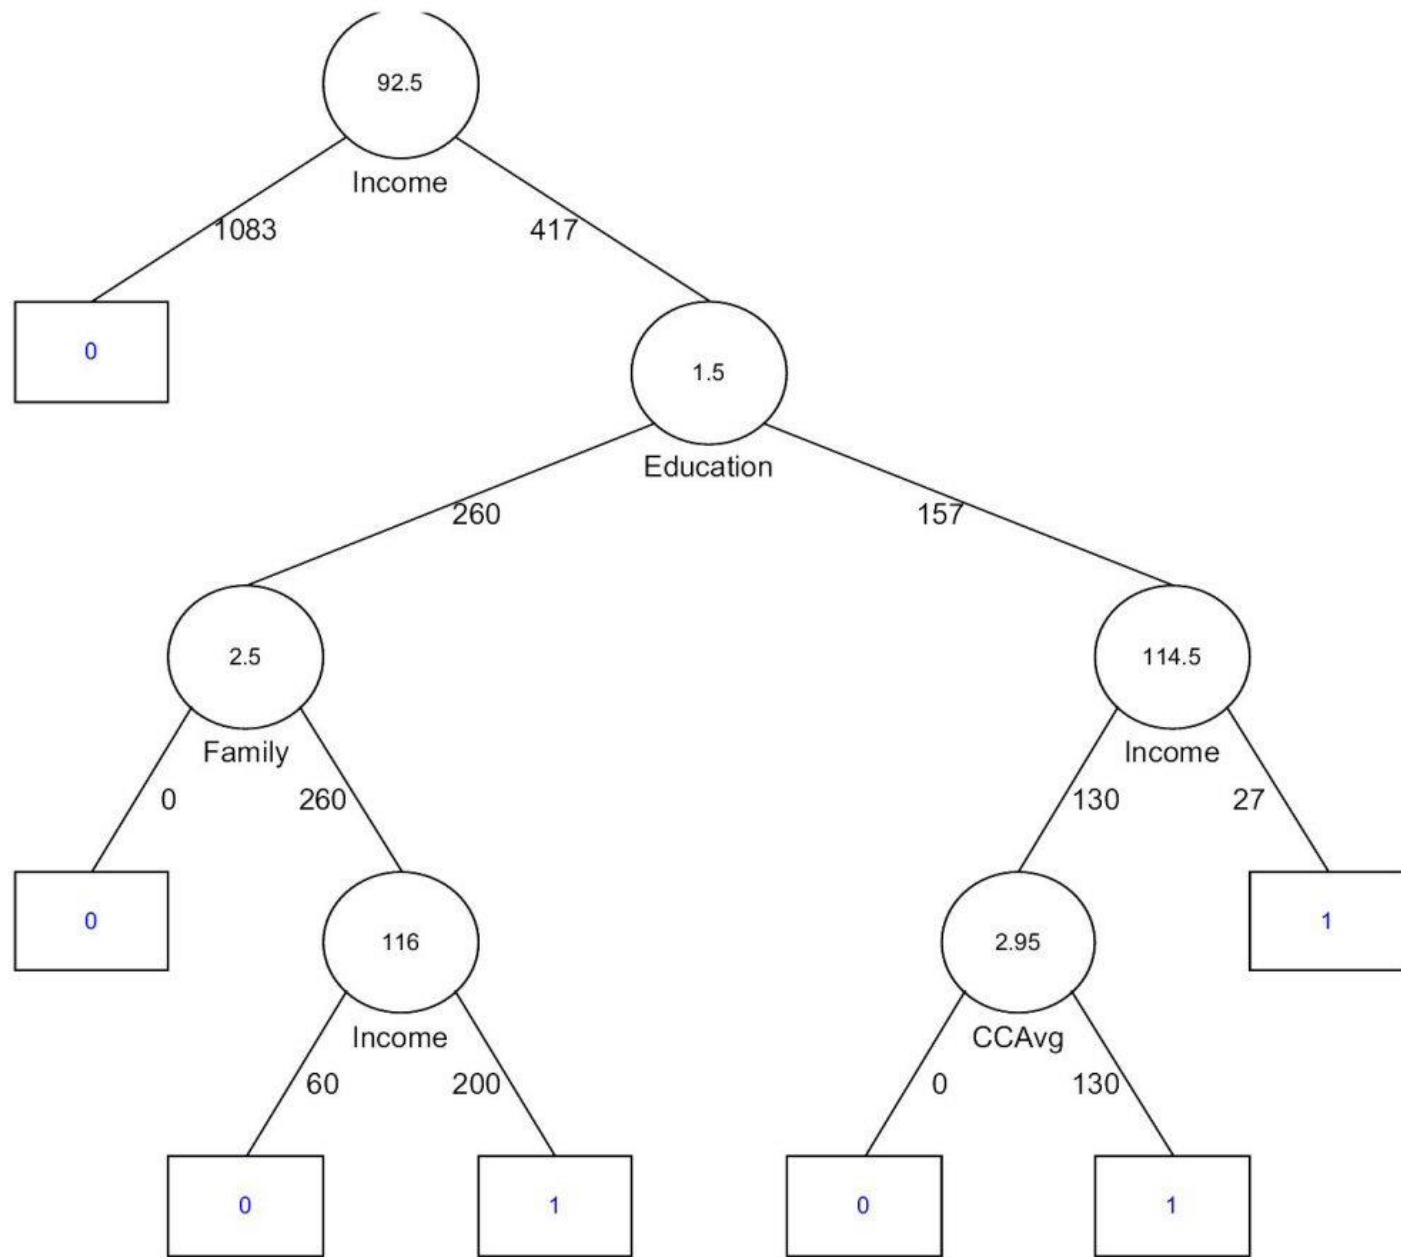
Example:

Goal: classify a record as “will accept credit card offer” or “will not accept”

Rule might be “IF (Income > 92.5) AND (Education < 1.5) AND (Family <= 2.5) THEN Class = 0 (nonacceptor)”

Also called CART, Decision Trees, or just Trees

Rules are represented by tree diagrams



Key Ideas

Recursive partitioning: Repeatedly split the records into two parts so as to achieve maximum homogeneity within the new parts

Pruning the tree: Simplify the tree by pruning peripheral branches to avoid overfitting

Recursive Partitioning Steps

Pick one of the predictor variables, x_i

Pick a value of x_i , say s_i , that divides the training data into two (not necessarily equal) portions

Measure how “pure” or homogeneous each of the resulting portions are
“Pure” = containing records of mostly one class

Algorithm tries different values of x_i and s_i to maximize purity in initial split

After you get a “maximum purity” split, repeat the process for a second split, and so on

Example: Riding Mowers

Goal: Classify 24 households as owning or not owning riding mowers

Predictors = Income, Lot Size

Income	Lot_Size	Ownership
60.0	18.4	owner
85.5	16.8	owner
64.8	21.6	owner
61.5	20.8	owner
87.0	23.6	owner
110.1	19.2	owner
108.0	17.6	owner
82.8	22.4	owner
69.0	20.0	owner
93.0	20.8	owner
51.0	22.0	owner
81.0	20.0	owner
75.0	19.6	non-owner
52.8	20.8	non-owner
64.8	17.2	non-owner
43.2	20.4	non-owner
84.0	17.6	non-owner
49.2	17.6	non-owner
59.4	16.0	non-owner
66.0	18.4	non-owner
47.4	16.4	non-owner
33.0	18.8	non-owner
51.0	14.0	non-owner
63.0	14.8	non-owner

How to split

Order records according to one variable, say lot size

Find midpoints between successive values

E.g. first midpoint is 14.4 (halfway between 14.0 and 14.8)

Divide records into those with lotsize > 14.4 and those < 14.4

After evaluating that split, try the next one, which is 15.4 (halfway between 14.8 and 16.0)

Note: Categorical Variables

Examine all possible ways in which the categories can be split.

E.g., categories A, B, C can be split 3 ways


{A} and {B, C}

{B} and {A, C}


{C} and {A, B}

With many categories, # of splits becomes huge


The first split: Lot Size = 19,000



Second Split: Income = \$84,000



After All Splits



Impurity measure: Gini Index

Gini Index for rectangle A containing m records

$$I(A) = 1 - \sum_{k=1}^m p_k^2$$

p = proportion of cases in rectangle A that belong to class k

- $I(A) = 0$ when all cases belong to same class
- Max value when all classes are equally represented (= 0.50 in binary case)

Impurity measure: Entropy

$$\textit{entropy}(A) = - \sum_{k=1}^m p_k \log_2(p_k)$$

p_k = proportion of cases (out of m) in rectangle A that belong to class k

Entropy ranges between 0 (most pure) and $\log_2(m)$ (equal representation of classes)

Impurity and Recursive Partitioning


Obtain overall impurity measure (weighted avg. of individual rectangles)

At each successive stage, compare this measure across all possible splits in all variables

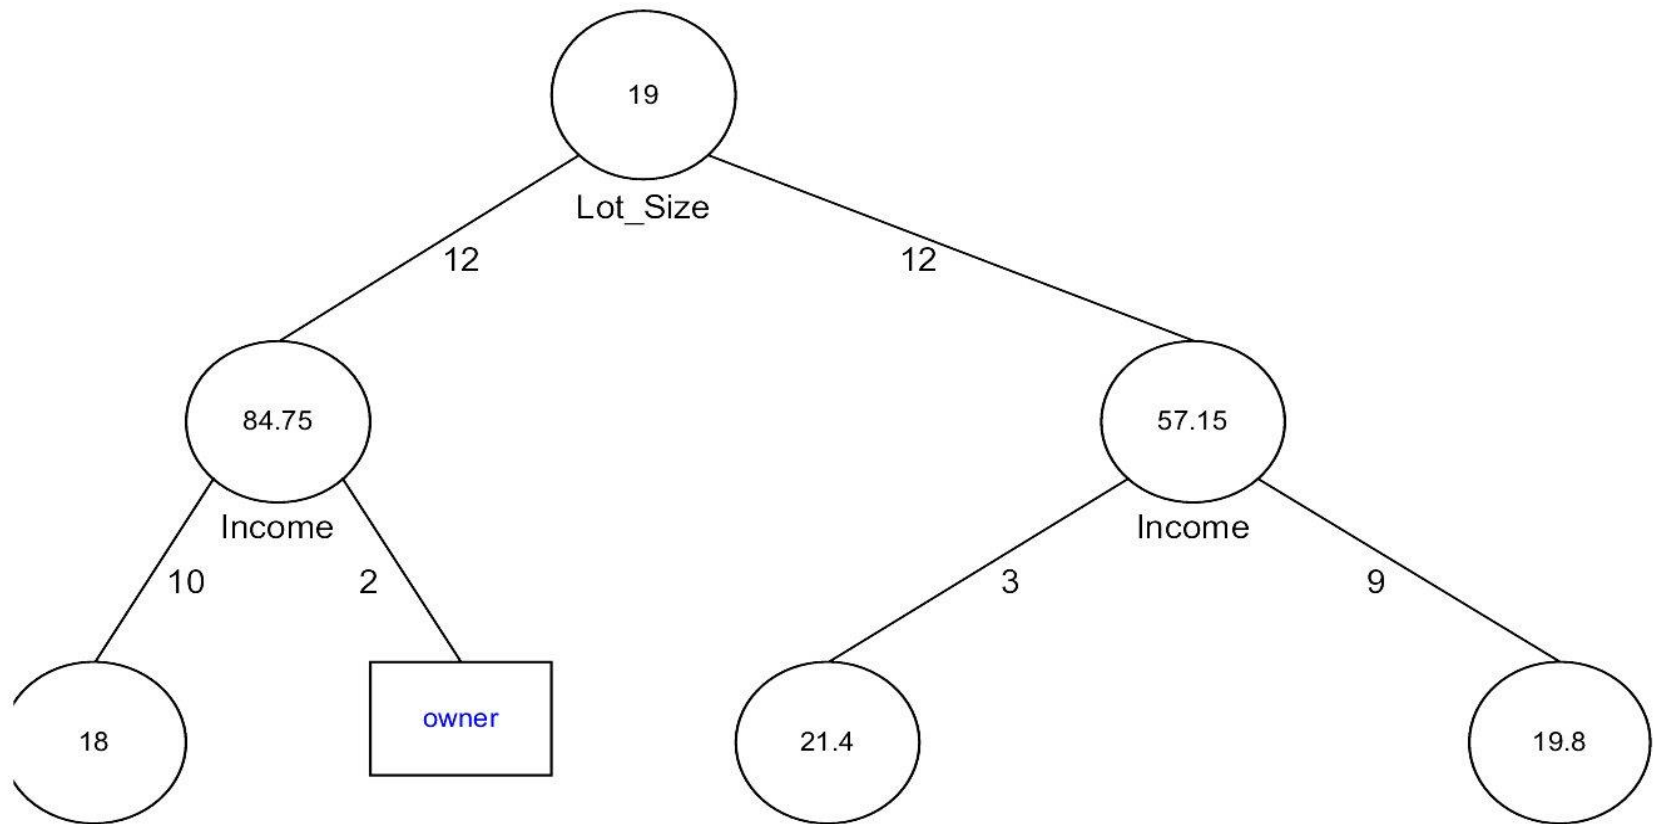
Choose the split that reduces impurity the most

Chosen split points become nodes on the tree

First Split – The Tree



Tree after three splits



Tree Structure

Split points become nodes on tree (circles with split value in center)

Rectangles represent “leaves” (terminal points, no further splits, classification value noted)

Numbers on lines between nodes indicate # cases

Read down tree to derive rule

E.g., If lot size < 19 , and if income > 84.75 , then class = “owner”

Determining Leaf Node Label

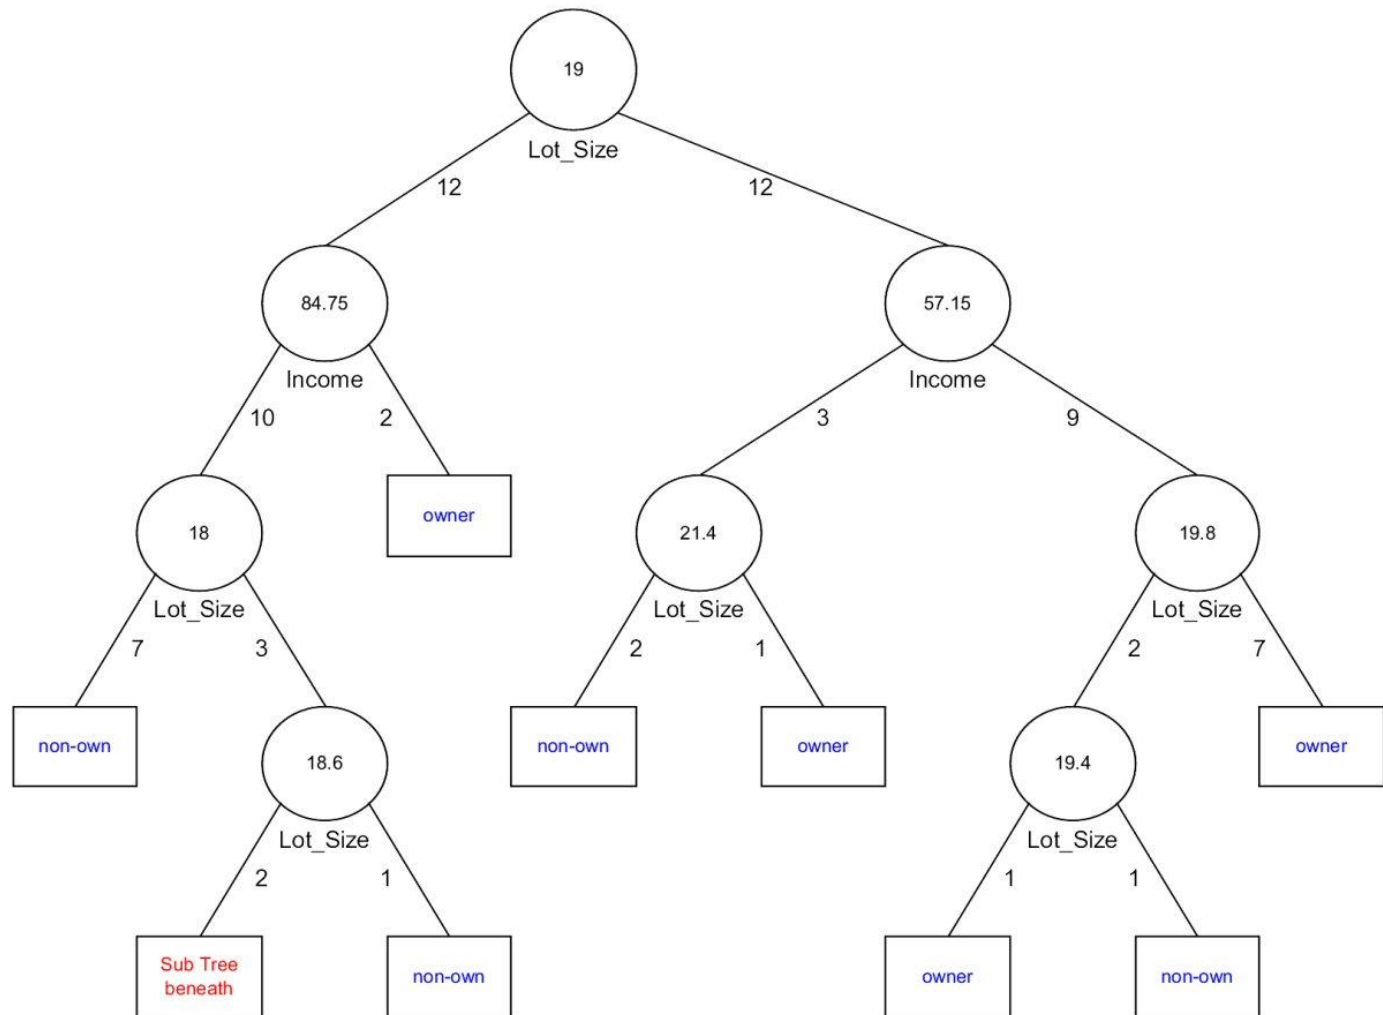
Each leaf node label is determined by “voting” of the records within it, and by the cutoff value

Records within each leaf node are from the training data

Default cutoff=0.5 means that the leaf node's label is the majority class.

Cutoff = 0.75: requires majority of 75% or more “1” records in the leaf to label it a “1” node

Tree after all splits



Preventing Overfitting: Stopping Tree Growth


Natural end of process is 100% purity in each leaf

This **overfits** the data, which end up fitting noise in the data

Overfitting leads to low predictive accuracy of new data

Past a certain point, the error rate for the validation data starts to increase

Full Tree Error Rate



CHAID

CHAID, older than CART, uses chi-square statistical test to limit tree growth

Splitting stops when purity improvement is not statistically significant

Pruning

CART lets tree grow to full extent, then prunes it back

Idea is to find that point at which the validation error begins to rise

Generate successively smaller trees by pruning leaves

At each pruning stage, multiple trees are possible

Use *cost complexity* to choose the best tree at that stage

Cost Complexity

$$CC(T) = Err(T) + \alpha L(T)$$

$CC(T)$ = cost complexity of a tree

$Err(T)$ = proportion of misclassified records

α = penalty factor attached to tree size (set by user)

Among trees of given size, choose the one with lowest CC

Do this for each size of tree

Using Validation Error to Prune

Pruning process yields a set of trees of different sizes and associated error rates

Two trees of interest:

Minimum error tree

Has lowest error rate on validation data

Best pruned tree

Smallest tree within one std. error of min. error

This adds a bonus for simplicity/parsimony

Regression Trees for Prediction

Used with continuous outcome variable

Procedure similar to classification tree

Many splits attempted, choose the one that minimizes impurity

Differences from CT

Prediction is computed as the **average** of numerical target variable in the rectangle (in CT it is majority vote)

Impurity measured by **sum of squared deviations** from leaf mean

Performance measured by RMSE (root mean squared error)

Advantages of trees

Easy to use, understand

Produce rules that are easy to interpret & implement

Variable selection & reduction is automatic

Do not require the assumptions of statistical models

Can work without extensive handling of missing data

Disadvantages

May not perform well where there is structure in the data that is not well captured by horizontal or vertical splits

Since the process deals with one variable at a time, no way to capture interactions between variables

Summary

Classification and Regression Trees are an easily understandable and transparent method for predicting or classifying new records

A tree is a graphical representation of a set of rules

Trees must be pruned to avoid over-fitting of the training data

As trees do not make any assumptions about the data structure, they usually require large samples